



ASSIGNMENT-4



APRIL 16, 2021

RISHABH SHARMA

20MAI0082

Github Link :- <https://github.com/rishabh5197/Deep-Learning-Assignments/tree/main/Assignment-4>

Task-1

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import tensorflow as tf
        4 import keras
        5 import matplotlib.pyplot as plt
        6 import seaborn as sns
        7 import scipy
        8 from sklearn.model_selection import train_test_split
```

```
In [2]: 1 (xtrain,ytrain),(xtest,ytest) = keras.datasets.cifar10.load_data()
        2 classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
In [3]: 1 plt.imshow(xtrain[0])
        2 plt.title(classes[ytrain[0][0]])
        3 plt.axis(False)
        4 plt.show()
```

frog



```
In [4]: 1 # selecting 50% less data from xtrain
2 (xtrain_new_50,xtest_new_50,ytrain_new_50,ytest_new_50) =train_test_split(xtrain,ytrain,test_size=0.5,random_state=
3 xtrain_1,xtest_1,ytrain_1,ytest_1 = train_test_split(xtrain_new_50,ytrain_new_50,test_size=0.3,random_state=0)
4 print("Xtrain original :- ",xtrain.shape)
5 print("Xtrain 50% selected from xtrain:-",xtrain_new_50.shape)
6 print("70% selected from that 50% training :- ",xtrain_1.shape)
```

Xtrain original :- (50000, 32, 32, 3)

Xtrain 50% selected from xtrain:- (25000, 32, 32, 3)

70% selected from that 50% training :- (17500, 32, 32, 3)

```
In [5]: 1 from keras.models import Sequential
2 from keras.layers import Conv2D,Activation,BatchNormalization,MaxPooling2D,Dense,Flatten,Dropout
```

```
In [6]: 1 from keras.utils import to_categorical
2
3 ytrain = to_categorical(ytrain,10)
4 ytrain_new_50 = to_categorical(ytrain_new_50,10)
5 ytrain_1 = to_categorical(ytrain_1,10)
6
7 ytest = to_categorical(ytest,10)
8 ytest_new_50 = to_categorical(ytest_new_50,10)
9 ytest_1 = to_categorical(ytest_1,10)
```

```
In [7]: 1 xtrain_1[0].shape
```

Out[7]: (32, 32, 3)

In [8]:

```
1 AlexNet = Sequential()
2
3 #1st Convolutional Layer
4 AlexNet.add(Conv2D(filters=96, input_shape=xtrain_1[0].shape, kernel_size=(11,11), strides=(4,4), padding='same'))
5 AlexNet.add(BatchNormalization())
6 AlexNet.add(Activation('relu'))
7
8 # 1st Maxpooling Layer
9 AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
10
11 #2nd Convolutional Layer
12 AlexNet.add(Conv2D(filters=256, kernel_size=(5, 5), strides=(1,1), padding='same'))
13 AlexNet.add(BatchNormalization())
14 AlexNet.add(Activation('relu'))
15
16 # 2nd Maxpooling Layer
17 AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
18
19 #3rd Convolutional Layer
20 AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
21 AlexNet.add(BatchNormalization())
22 AlexNet.add(Activation('relu'))
23
24 #4th Convolutional Layer
25 AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
26 AlexNet.add(BatchNormalization())
27 AlexNet.add(Activation('relu'))
28
29 #5th Convolutional Layer
30 AlexNet.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same'))
31 AlexNet.add(BatchNormalization())
32 AlexNet.add(Activation('relu'))
33
34 # 3rd Maxpooling Layer
35 AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
36
37 #Passing it to a Fully Connected layer
38 AlexNet.add(Flatten())
39
40 # 1st Fully Connected Layer
41 AlexNet.add(Dense(4096, input_shape=(32,32,3,)))
42 AlexNet.add(BatchNormalization())
```

```

43 AlexNet.add(Activation('relu'))
44
45 # Add Dropout to prevent overfitting
46 AlexNet.add(Dropout(0.4))
47
48 #2nd Fully Connected Layer
49 AlexNet.add(Dense(4096))
50 AlexNet.add(BatchNormalization())
51 AlexNet.add(Activation('relu'))
52 #Add Dropout
53 AlexNet.add(Dropout(0.4))
54
55 #3rd Fully Connected Layer
56 # AlexNet.add(Dense(1000)) # this is being removed to perform ConvNet as fixed feature extractor
57 AlexNet.add(BatchNormalization())
58 AlexNet.add(Activation('relu'))
59 #Add Dropout
60 AlexNet.add(Dropout(0.4))
61
62 #Output Layer
63 AlexNet.add(Dense(10))
64 AlexNet.add(BatchNormalization())
65 AlexNet.add(Activation('softmax'))
66
67 #Model Summary
68 AlexNet.summary()
69 # https://cs231n.github.io/transfer-learning/

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 8, 8, 96)	34944
batch_normalization (Batch Normalization)	(None, 8, 8, 96)	384
activation (Activation)	(None, 8, 8, 96)	0
max_pooling2d (MaxPooling2D)	(None, 4, 4, 96)	0
conv2d_1 (Conv2D)	(None, 4, 4, 256)	614656
batch_normalization_1 (Batch Normalization)	(None, 4, 4, 256)	1024

activation_1 (Activation)	(None, 4, 4, 256)	0
max_pooling2d_1 (MaxPooling2)	(None, 2, 2, 256)	0
conv2d_2 (Conv2D)	(None, 2, 2, 384)	885120
batch_normalization_2 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_2 (Activation)	(None, 2, 2, 384)	0
conv2d_3 (Conv2D)	(None, 2, 2, 384)	1327488
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_3 (Activation)	(None, 2, 2, 384)	0
conv2d_4 (Conv2D)	(None, 2, 2, 256)	884992
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 256)	1024
activation_4 (Activation)	(None, 2, 2, 256)	0
max_pooling2d_2 (MaxPooling2)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 4096)	1052672
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
activation_5 (Activation)	(None, 4096)	0
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
activation_6 (Activation)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0

batch_normalization_7 (Batch Normalization)	(None, 4096)	16384
activation_7 (Activation)	(None, 4096)	0
dropout_2 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 10)	40970
batch_normalization_8 (Batch Normalization)	(None, 10)	40
activation_8 (Activation)	(None, 10)	0
=====		
Total params: 21,676,850		
Trainable params: 21,649,502		
Non-trainable params: 27,348		

In [9]: 1 AlexNet.compile(loss = 'categorical_crossentropy', optimizer= 'adam', metrics=['accuracy'])

In [10]: 1 # checkpoint = ModelCheckpoint('AlexNet.h5',save_best_only=True, monitor='val_accuracy',mode='max')
2 history = AlexNet.fit(xtrain_1,ytrain_1,epochs=100,batch_size=32)

```
547/547 [=====] - 7s 13ms/step - loss: 0.1324 - accuracy: 0.9609
Epoch 32/100
547/547 [=====] - 7s 13ms/step - loss: 0.1261 - accuracy: 0.9617
Epoch 33/100
547/547 [=====] - 7s 13ms/step - loss: 0.1211 - accuracy: 0.9665
Epoch 34/100
547/547 [=====] - 7s 13ms/step - loss: 0.1212 - accuracy: 0.9645
Epoch 35/100
547/547 [=====] - 7s 13ms/step - loss: 0.1023 - accuracy: 0.9682
Epoch 36/100
547/547 [=====] - 7s 13ms/step - loss: 0.1170 - accuracy: 0.9663
Epoch 37/100
547/547 [=====] - 7s 13ms/step - loss: 0.0979 - accuracy: 0.9706
Epoch 38/100
547/547 [=====] - 7s 13ms/step - loss: 0.1012 - accuracy: 0.9693
Epoch 39/100
547/547 [=====] - 7s 13ms/step - loss: 0.0876 - accuracy: 0.9740
Epoch 40/100
547/547 [=====] - 7s 13ms/step - loss: 0.1026 - accuracy: 0.9692
Epoch 41/100
```


In [11]: 1 history = pd.DataFrame(history.history)

In [12]: 1 history


Out[12]:

	loss	accuracy
0	1.749619	0.361543
1	1.536217	0.454286
2	1.422057	0.498286
3	1.329333	0.536057
4	1.245421	0.559543
...
95	0.032239	0.990114
96	0.031350	0.990800
97	0.031872	0.989829
98	0.027423	0.992000
99	0.032049	0.990057

100 rows × 2 columns

```
In [13]: 1 history.plot.line(figsize=(16,10),)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f539a459250>
```



```
In [14]: 1 ypred = np.argmax(AlexNet.predict(xtest_1),axis=1)
          2
          3 ytrain = np.argmax(ytrain,axis=1)
          4 ytrain_new_50 = np.argmax(ytrain_new_50,axis=1)
          5 ytrain_1 = np.argmax(ytrain_1,axis=1)
          6
          7 ytest = np.argmax(ytest,axis=1)
          8 ytest_new_50 = np.argmax(ytest_new_50,axis=1)
          9 ytest_1 = np.argmax(ytest_1,axis=1)
```

```
In [15]: 1 from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
          2 accuracy_score(ypred,ytest_1)
```

Out[15]: 0.5650666666666667

```
In [16]: 1 ytrain = to_categorical(ytrain,10)
          2 ytrain_new_50 = to_categorical(ytrain_new_50,10)
          3 ytrain_1 = to_categorical(ytrain_1,10)
          4
          5 ytest = to_categorical(ytest,10)
          6 ytest_new_50 = to_categorical(ytest_new_50,10)
          7 ytest_1 = to_categorical(ytest_1,10)
```

In [17]: 1 history = AlexNet.fit(xtrain,ytrain,epochs=100,batch_size=32)

```
Epoch 1/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.8553 - accuracy: 0.7369
Epoch 2/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.6615 - accuracy: 0.7933
Epoch 3/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.5335 - accuracy: 0.8344
Epoch 4/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.4304 - accuracy: 0.8664
Epoch 5/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.3458 - accuracy: 0.8939
Epoch 6/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.2831 - accuracy: 0.9126
Epoch 7/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.2330 - accuracy: 0.9264
Epoch 8/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.1898 - accuracy: 0.9396
Epoch 9/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.1634 - accuracy: 0.9495
Epoch 10/100
1563/1563 [=====] - 20s 13ms/step - loss: 0.1422 - accuracy: 0.9547
```

In [18]: 1 history = pd.DataFrame(history.history)

In [19]:

1 history


Out[19]:

	loss	accuracy
0	0.855332	0.73690
1	0.661502	0.79328
2	0.533495	0.83442
3	0.430364	0.86640
4	0.345819	0.89390
...
95	0.022198	0.99338
96	0.018639	0.99414
97	0.019822	0.99392
98	0.018850	0.99404
99	0.018977	0.99400

100 rows × 2 columns

In [20]: 1 history.plot.line(figsize=(16,10))

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f53d01a0a90>



```
In [21]: 1 ypred = np.argmax(AlexNet.predict(xtest),axis=1)
          2
          3 ytrain = np.argmax(ytrain,axis=1)
          4 ytrain_new_50 = np.argmax(ytrain_new_50,axis=1)
          5 ytrain_1 = np.argmax(ytrain_1,axis=1)
          6
          7 ytest = np.argmax(ytest,axis=1)
          8 ytest_new_50 = np.argmax(ytest_new_50,axis=1)
          9 ytest_1 = np.argmax(ytest_1,axis=1)
```

```
In [22]: 1 from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
          2 accuracy_score(ypred,ytest)
```

Out[22]: 0.6501

Task-2


```
In [1]: 1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pandas as pd
4 import seaborn as sns
5 import keras as k
6 import tensorflow as tf
7 from keras.datasets import cifar10
8 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
9 from sklearn.model_selection import train_test_split
```

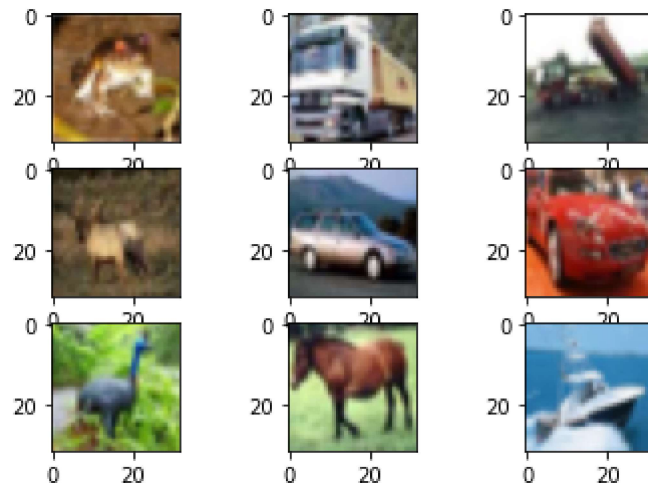
```
In [2]: 1 (trainX, trainy), (testX, testy) = cifar10.load_data()
2 print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
3 print('Test: X=%s, y=%s' % (testX.shape, testy.shape))
4 for i in range(9):
5     plt.subplot(330 + 1 + i)
6     plt.imshow(trainX[i])
7     plt.show()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> (<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>)

170500096/170498071 [=====] - 2s 0us/step

Train: X=(50000, 32, 32, 3), y=(50000, 1)

Test: X=(10000, 32, 32, 3), y=(10000, 1)



```
In [3]: 1 from keras.utils import to_categorical  
        2 trainy = to_categorical(trainy)  
        3 testy = to_categorical(testy)
```

```
In [4]: 1 from keras.layers import Conv2D,MaxPooling2D,Flatten,Dense,BatchNormalization  
        2 from keras.models import Sequential  
        3 from keras.optimizers import SGD  
        4 from keras.preprocessing.image import ImageDataGenerator
```

```
In [5]: 1 datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
```

```
In [6]: 1 model = Sequential()
2 model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(32, 32, 3)))
3 model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
4 model.add(MaxPooling2D((2, 2)))
5 model.add(BatchNormalization()) # Adding Batch Normalization
6 model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
7 model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
8 model.add(MaxPooling2D((2, 2)))
9 model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
10 model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
11 model.add(MaxPooling2D((2, 2)))
12 model.add(Flatten())
13 model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
14 model.add(Dense(10, activation='softmax'))
15 # compile model
16 opt = SGD(lr=0.001, momentum=0.9)
17 model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
18 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
batch_normalization (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0

flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 550,698		
Trainable params: 550,634		
Non-trainable params: 64		


In [7]: 1 it_train = datagen.flow(trainX, trainy, batch_size=64)

In [8]: 1 steps = int(trainX.shape[0] / 64)
2 history = model.fit_generator(it_train, steps_per_epoch=steps, epochs=20, validation_data=(testX, testy), verbose=0)

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
warnings.warn("`Model.fit_generator` is deprecated and "

In [9]: 1 history = pd.DataFrame(history.history)

In [16]: 1 history.plot.line(figsize=(16,10)).grid("whitegrid")



```
In [11]: 1 ypred = np.argmax(model.predict(testX),axis=1)
```

```
In [12]: 1 ypred
```

```
Out[12]: array([3, 8, 8, ..., 5, 1, 7])
```

```
In [13]: 1 testty = np.argmax(testy,axis=1)
```

```
In [14]: 1 testty
```

```
Out[14]: array([3, 8, 8, ..., 5, 1, 7])
```

```
In [15]: 1 accuracy_score(ypred,testty)
```

```
Out[15]: 0.7764
```