

```
In [1]: 1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pandas as pd
4 import seaborn as sns
5 import keras as k
6 import tensorflow as tf
7 from keras.datasets import cifar10
8 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
9 from sklearn.model_selection import train_test_split
```

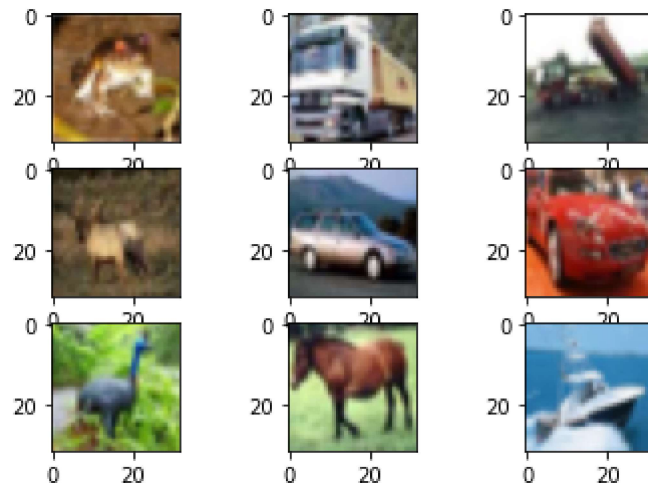
```
In [2]: 1 (trainX, trainy), (testX, testy) = cifar10.load_data()
2 print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
3 print('Test: X=%s, y=%s' % (testX.shape, testy.shape))
4 for i in range(9):
5     plt.subplot(330 + 1 + i)
6     plt.imshow(trainX[i])
7     plt.show()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> (<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>)

170500096/170498071 [=====] - 2s 0us/step

Train: X=(50000, 32, 32, 3), y=(50000, 1)

Test: X=(10000, 32, 32, 3), y=(10000, 1)



```
In [3]: 1 from keras.utils import to_categorical  
2 trainy = to_categorical(trainy)  
3 testy = to_categorical(testy)
```

```
In [4]: 1 from keras.layers import Conv2D,MaxPooling2D,Flatten,Dense,BatchNormalization  
2 from keras.models import Sequential  
3 from keras.optimizers import SGD  
4 from keras.preprocessing.image import ImageDataGenerator
```

```
In [5]: 1 datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
```

```
In [6]: 1 model = Sequential()
2 model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(32, 32, 3)))
3 model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
4 model.add(MaxPooling2D((2, 2)))
5 model.add(BatchNormalization()) # Adding Batch Normalization
6 model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
7 model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
8 model.add(MaxPooling2D((2, 2)))
9 model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
10 model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
11 model.add(MaxPooling2D((2, 2)))
12 model.add(Flatten())
13 model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
14 model.add(Dense(10, activation='softmax'))
15 # compile model
16 opt = SGD(lr=0.001, momentum=0.9)
17 model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
18 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
batch_normalization (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0

flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 550,698		
Trainable params: 550,634		
Non-trainable params: 64		


In [7]: 1 it_train = datagen.flow(trainX, trainy, batch_size=64)

In [8]: 1 steps = int(trainX.shape[0] / 64)
2 history = model.fit_generator(it_train, steps_per_epoch=steps, epochs=20, validation_data=(testX, testy), verbose=0)

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
warnings.warn("`Model.fit_generator` is deprecated and "

In [9]: 1 history = pd.DataFrame(history.history)

In [16]: 1 history.plot.line(figsize=(16,10)).grid("whitegrid")



```
In [11]: 1 ypred = np.argmax(model.predict(testX),axis=1)
```

```
In [12]: 1 ypred
```

```
Out[12]: array([3, 8, 8, ..., 5, 1, 7])
```

```
In [13]: 1 testty = np.argmax(testy,axis=1)
```

```
In [14]: 1 testty
```

```
Out[14]: array([3, 8, 8, ..., 5, 1, 7])
```

```
In [15]: 1 accuracy_score(ypred,testty)
```

```
Out[15]: 0.7764
```